



From automatic
differentiation to message
passing

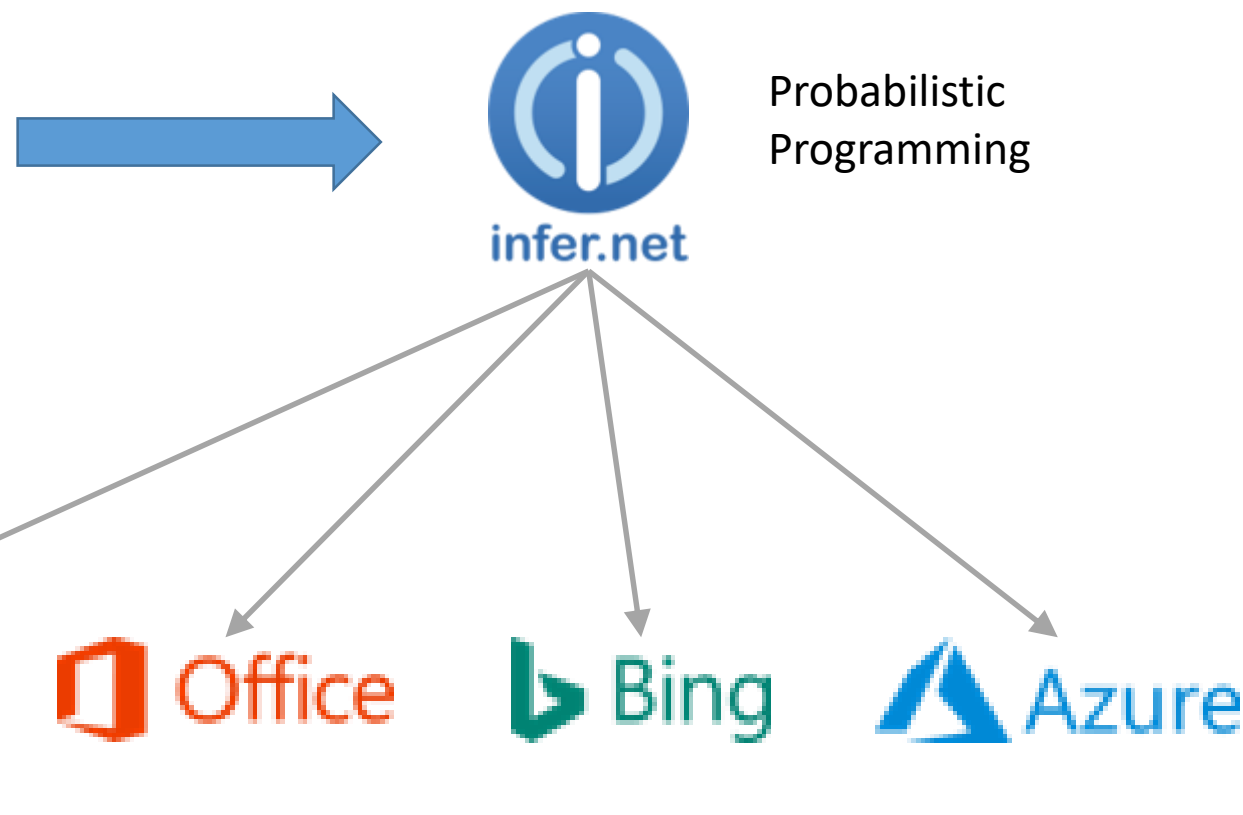
Tom Minka
Microsoft Research

What I do



Algorithms for probabilistic inference

- Expectation Propagation
- Non-conjugate variational message passing
- A* sampling



Machine Learning Language



- A machine learning language should (among other things) simplify implementation of machine learning algorithms

Machine Learning Language



- A **general-purpose** machine learning language should (among other things) simplify implementation of **all** machine learning algorithms

Roadmap



1. Automatic Differentiation
2. AutoDiff lacks approximation
3. Message passing generalizes AutoDiff
4. Compiling to message passing



1. Automatic / algorithmic differentiation

Recommended reading



- “Evaluating derivatives” by Griewank and Walther (2008)

Programs are the new formulas



- Programs can specify mathematical functions more compactly than formulas
- Program is not a black box: undergoes analysis and transformation
- Numbers are assumed to have infinite precision

Multiply-all example



- As formulas:
- $f = \prod_i x_i$
- $df = \sum_i dx_i \prod_{j \neq i} x_j$

Multiply-all example



Input program

```
c[1] = x[1]
for i = 2 to n
  c[i] = c[i-1]*x[i]
f = c[n]
```



Derivative program

```
dc[1] = dx[1]
for i = 2 to n
  dc[i] = dc[i-1]*x[i] + c[i-1]*dx[i]
df = dc[n]
```

$$f = \prod_i x_i$$

$$df = \sum_i dx_i \prod_{j \neq i} x_j$$

Phases of AD

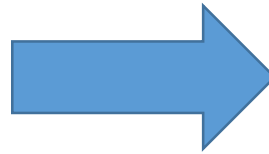
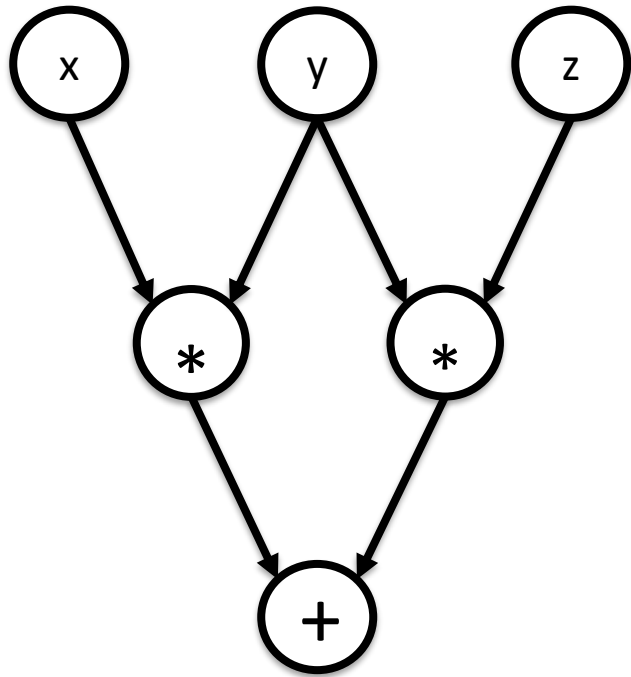


- Execution
 - Replace every operation with a linear one
- Accumulation
 - Collect linear coefficients

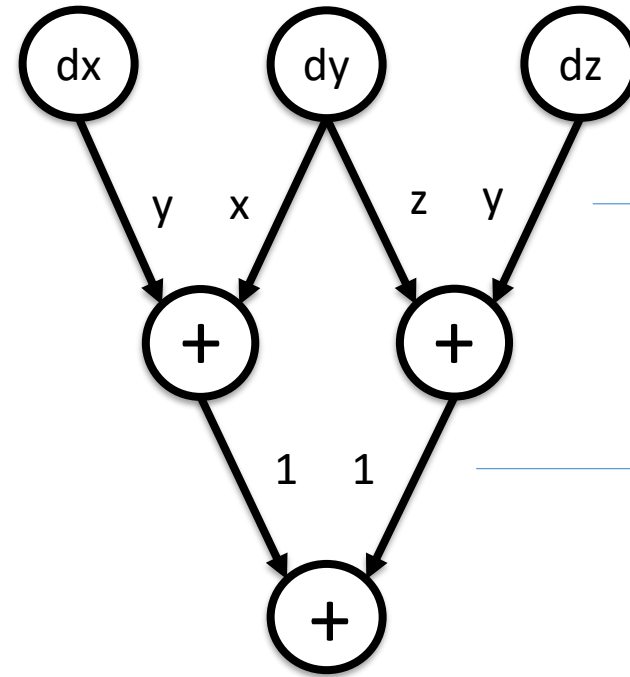
Execution phase



$$x*y + y*z$$



$$dx*y + x*dy + dy*z + y*dz$$

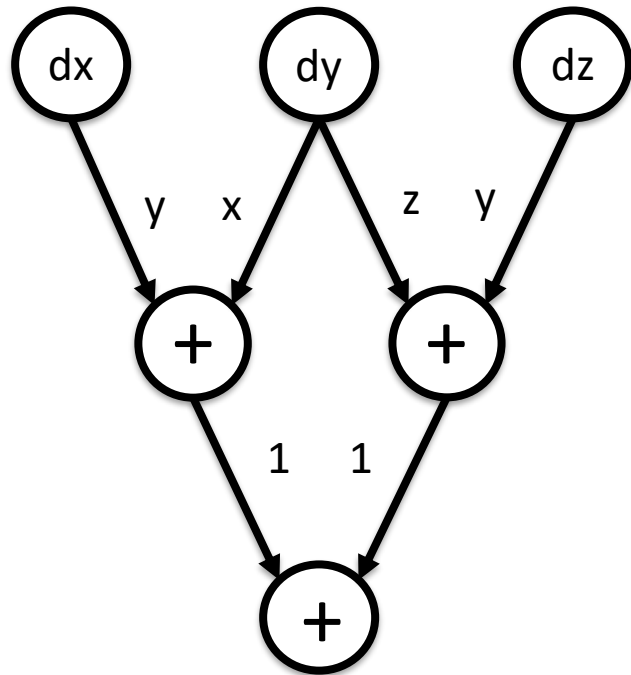


Scale
factors

Accumulation phase



$$dx*y + x*dy + dy*z + y*dz \quad (\text{Forward})$$



(Reverse)

coefficient of $dx = 1*y$

coefficient of $dy = 1*x + 1*z$

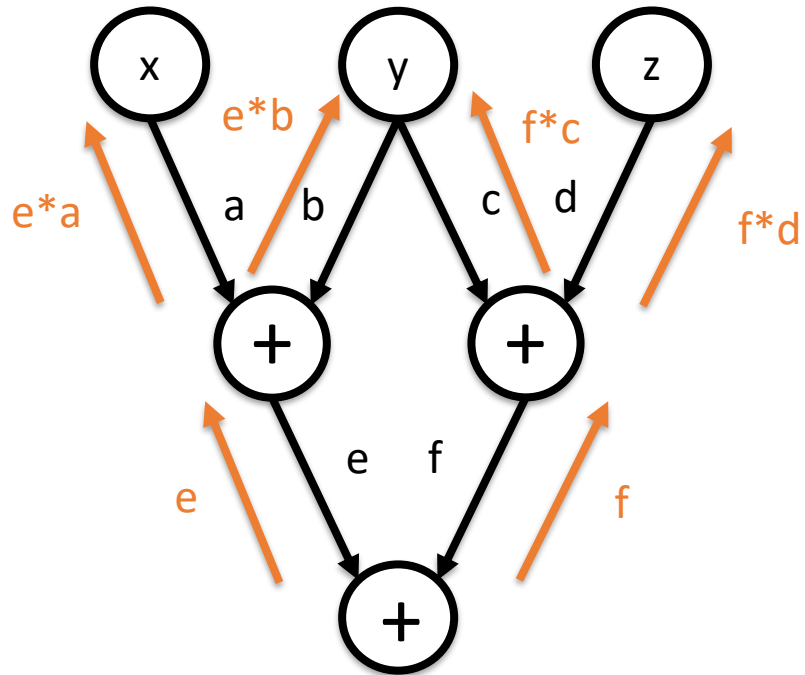
coefficient of $dz = 1*y$

Gradient vector = $(1*y, 1*x + 1*z, 1*y)$

Linear composition



$$e*(a*x + b*y) + f*(c*y + d*z)$$



$$(e*a)*x +$$


$$(e*b + f*c)*y +$$

$$(f*d)*z$$

Dynamic programming



- Reverse accumulation is dynamic programming
- Backward message is sum over paths to output



Source-to-source translation



- Tracing approach builds a graph during execution phase, then accumulates it
- Source-to-source produces a gradient program matching structure of original

Multiply-all example



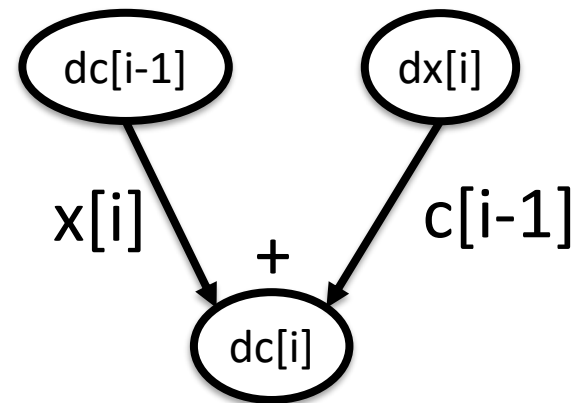
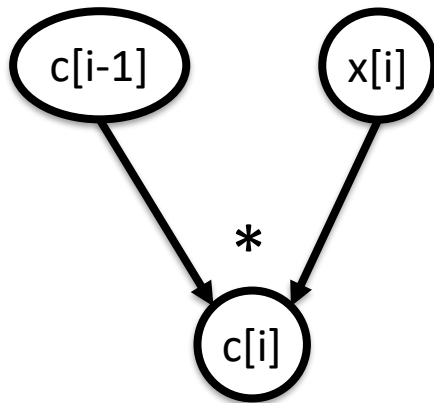
Input program

```
c[1] = x[1]
for i = 2 to n
  c[i] = c[i-1]*x[i]
return c[n]
```



Derivative program

```
dc[1] = dx[1]
for i = 2 to n
  dc[i] = dc[i-1]*x[i] + c[i-1]*dx[i]
return dc[n]
```



Multiply-all example




Derivative program

```
dc[1] = dx[1]
for i = 2 to n
  dc[i] = dc[i-1]*x[i] + c[i-1]*dx[i]
return dc[n]
```

Gradient program

```
dcB[n] = 1
for i = n downto 2
  dcB[i-1] = dcB[i]*x[i]
  dxB[i] = dcB[i]*c[i-1]
dxB[1] = dcB[1]
return dxB
```



General case



$$c = f(x,y)$$




$$dc = df1(x,y) * dx + df2(x,y) * dy$$



$$dx_B = dc_B * df1(x,y)$$

$$dy_B = dc_B * df2(x,y)$$



Fan-out




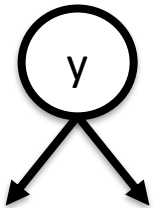
- If a variable is read multiple times, we need to add its backward messages
- Non-incremental approach:
transform program so that each variable is defined and used at most once on every execution path

Fan-out example




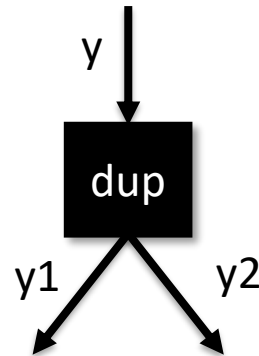
Input program

```
a = x * y
b = y * z
c = a + b
```



Edge program

```
(y1,y2) = dup(y)
a = x * y1
b = y2 * z
c = a + b
```



Gradient program

```
aB = cB
bB = cB
y2B = bB * z
y1B = aB * x
yB = y1B + y2B
...
```

Summary of AutoDiff



	AD	Message passing
Programs not formulas	Yes	Yes
Graph structure / sparsity	Yes	Yes
Source-to-source	Yes	Yes
Only one execution path	Yes	Not always
Single forward-backward sweep	Yes	Not always
Exact	Yes	Not always



2. AutoDiff lacks approximation

Approximate gradients for big models



- Mini-batching
- User changes input program to be approximate, then computes exact gradient

$$\begin{aligned} \nabla \sum_{i=1}^n f_i(\theta) &\approx \\ \nabla \frac{n}{m} \sum_{s \sim (1:n)} f_s(\theta) &= \\ \frac{n}{m} \sum_{s \sim (1:n)} \nabla f_s(\theta) &\quad (\text{AutoDiff}) \end{aligned}$$

Black-box variational inference



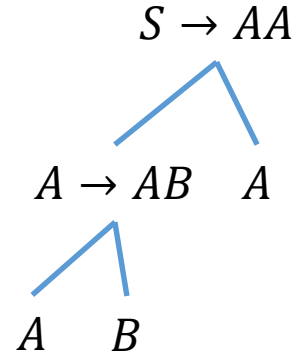
1. Approximate the marginal log-likelihood with a lower bound
2. Approximate the lower bound by importance sampling
3. Compute exact gradient of approximation

$$\int p(x, D) dx \geq -KL(q \parallel p)$$

AutoDiff in Tractable Models



- AutoDiff can mechanically derive reverse summation algorithms for tractable models
 - Markov chains, Bayesian networks (Darwiche, 2003)
 - Generative grammars, Parse trees (Eisner, 2016)
- Posterior expectations are derivatives of marginal log-likelihood, which can be computed exactly
 - User must provide forward summation algorithm



Approximation in Tractable Models



- Approximation is useful in tractable models
 - Sparse forward-backward (Pal et al, 2006)
 - Beam parsing (Goodman, 1997)
- Cannot be obtained through AutoDiff of an approximate model
- Neither can Viterbi

MLL should facilitate approximations



- Expectations
- Fixed-point iteration
 - Optimization
 - Root finding
- Should all be natively supported



3. Message-passing generalizes autodiff

Message-passing



- Approximate reasoning about exponential state space of a program, along all execution paths
- Propagates state summaries in both directions
- Forward can depend on backward and vice versa
- Iterate to convergence

Interval constraint propagation

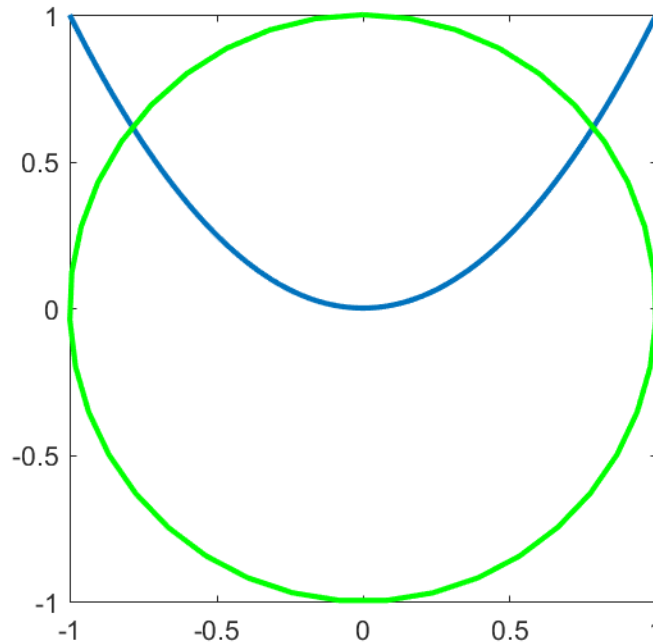


- What is largest and smallest value each variable could have?
- Each operation in program is interpreted as a constraint between inputs and output
- Propagates information forward and backward until convergence

Circle-parabola example



Find (x, y) that satisfies $x^2 + y^2 = 1$
and $y = x^2$

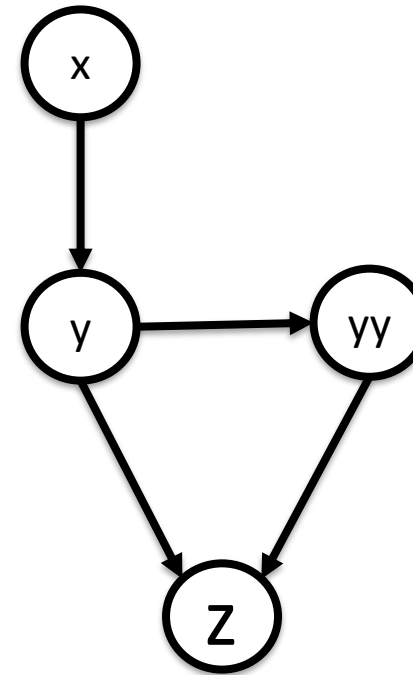


Circle-parabola program



Input program

```
y = x^2  
yy = y^2  
z = y + yy  
assert(z == 1)
```



Interval propagation program



Input program

$y = x^2$

$yy = y^2$

$z = y + yy$

$\text{assert}(z == 1)$



Edge program

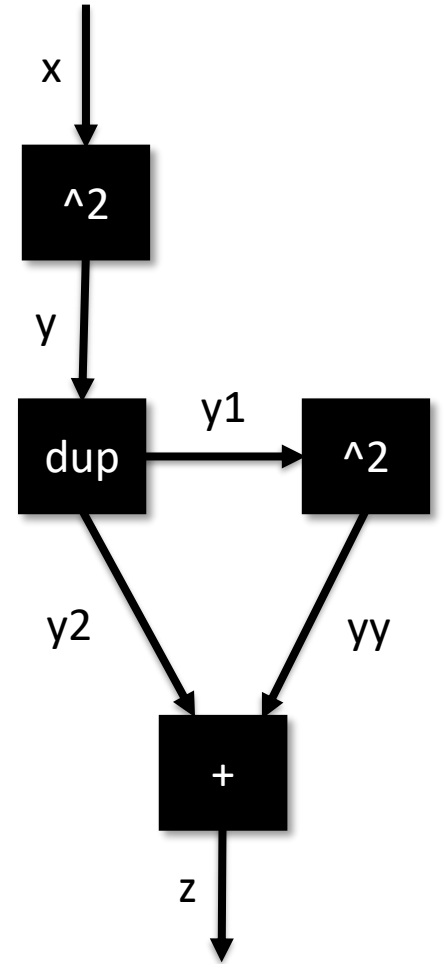
$y = x^2$

$(y1, y2) = \text{dup}(y)$

$yy = y1^2$

$z = y2 + yy$

$\text{assert}(z == 1)$



Interval propagation program



Edge program

$y = x^2$

$(y1, y2) = \text{dup}(y)$

$yy = y1^2$

$z = y2 + yy$

$\text{assert}(z == 1)$

Message program

Until convergence:

$yF = xF^2$

$y1F = yF \cap y2B$

$y2F = yF \cap y1B$


$yyF = y1F^2$

$y1B = \text{sqrt}(y1F, yyB)$

$y2B = zB - yyF$

$yyB = zB - y2F$

$zB = [1, 1]$



Running $\wedge 2$ backwards



$$yy = y1^2 \quad \longrightarrow \quad y1B = \text{sqrt}(y1F, yyB) \\ = \text{project}[y1F \cap \text{sqrt}(yyB)]$$

$$yyB = [2, 4]$$

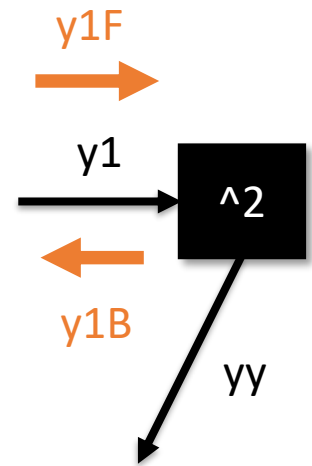
$$\text{sqrt}(yyB) = [-2, -1] \cup [1, 2]$$

$$y1F = [0, 10]$$

$$y1F \cap \text{sqrt}(yyB) = [] \cup [1, 2]$$

$$\text{project}[y1F \cap \text{sqrt}(yyB)] = [1, 2]$$

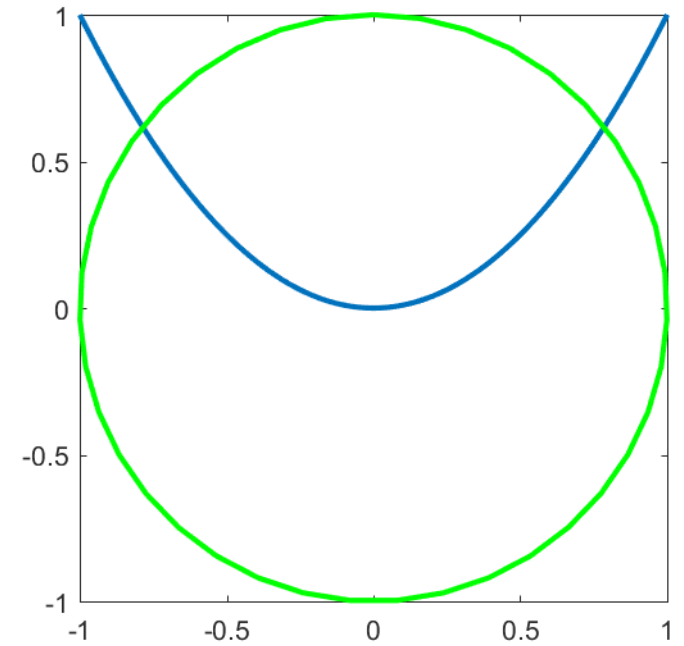
$$y1F \cap \text{project}[\text{sqrt}(yyB)] = [0, 2]$$



Results



- If all intervals start $(-\infty, \infty)$ then $x \rightarrow (-1, 1)$ (overestimate)
- Apply subdivision
- Starting at $x = (0.1, 1)$ gives $x \rightarrow (0.786, 0.786)$



Interval propagation program



Until convergence:

$$yF = xF^2$$

$$xB = \text{sqrt}(xF, yB)$$

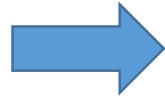
$$yB = y1B \cap y2B$$

$$y1F = yF \cap y2B$$

$$y2F = yF \cap y1B$$

...

$$zB = [1,1]$$



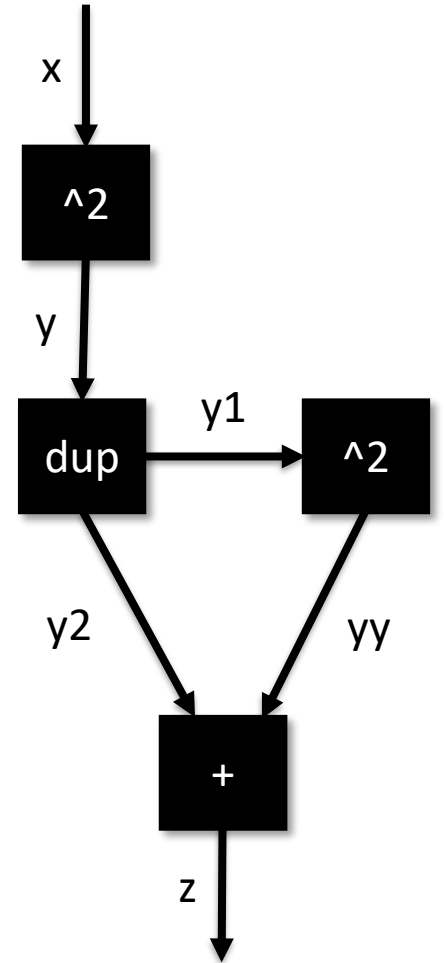
$$yF = xF^2$$

$$zB = [1,1]$$

Until convergence:
(perform updates)

$$yB = y1B \cap y2B$$

$$xB = \text{sqrt}(xF, yB)$$

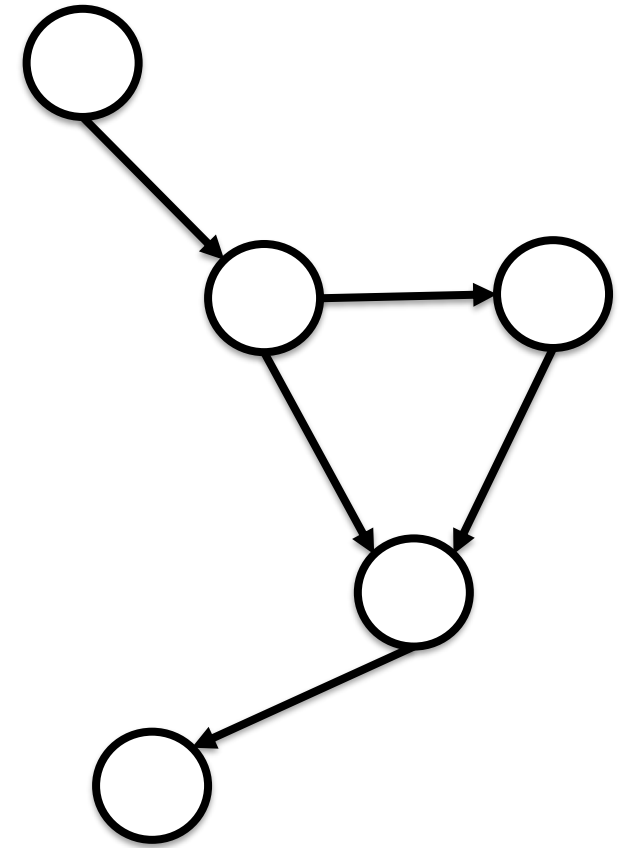


Typical message-passing program



1. Pass messages into the loopy core
2. Iterate
3. Pass messages out of the loopy core

Analogous to Stan's "transformed data" and "generated quantities"



Simplifications of message-passing



- Message dependencies dictate execution
- If forward messages do not depend on backward, becomes non-iterative
- If forward messages only include single state, only one control path is explored
- AutoDiff has both properties



Other message-passing algorithms

Probabilistic Programming



- Probabilistic programs are the new Bayesian networks
- Using a program to specify a probabilistic model
- Program is not a black box: undergoes analysis and transformation to help inference

Loopy belief propagation




- Loopy belief propagation has same structure as interval propagation, but using distributions
 - Gives forward and backward summations for tractable models
- Expectation propagation adds projection steps
 - Approximate expectations for intractable models
 - Parameter estimation in non-conjugate models

Gradient descent



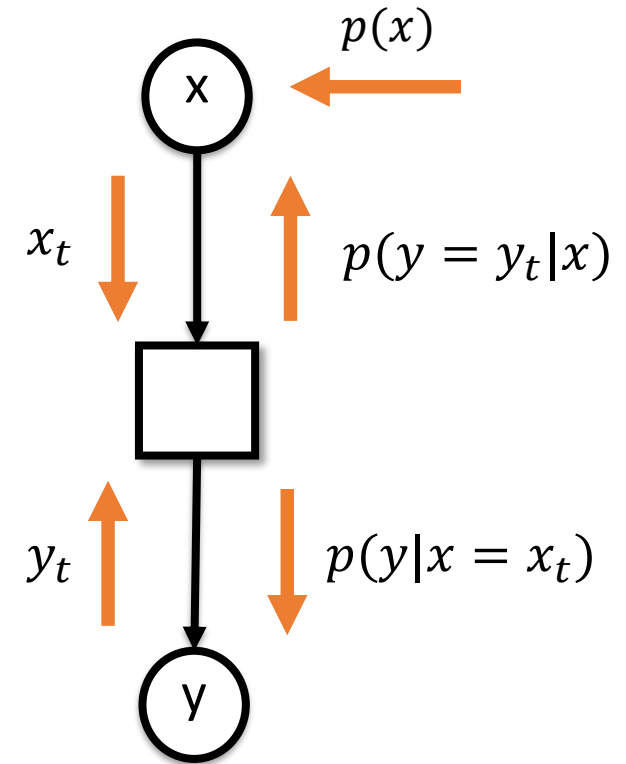
- Parameters send current value out, receive gradients in, take a step
 - Gradients fall out of EP equations
- Part of the same iteration loop



Gibbs sampling



- Variables send current value out, receive conditional distributions in
- Collapsed variables send/receive distributions as in BP
 - No need to collapse in the model





Thanks!

Model-based machine learning book: <http://mbmlbook.com/>

Infer.NET is open source: <http://dotnet.github.io/infer>